

Deliverable report

Antonio Gelain: Mat: 258784,
antonio.gelain@studenti.unitn.it,
antonio.gelain@eagletrt.it,
GitRepo: <https://github.com/Tonidotpy/GPU-Computing-2025-258784>

Abstract—Sparse Matrix-Vector multiplication (SpMV) is a fundamental operation in fields such as scientific computing, graph analysis and machine learning, which can often become a performance bottleneck due to the large amount of data which needs to be processed. This report will cover the results of various implementations and optimization strategies used to achieve faster computation leveraging both the Graphics Processing Units (GPUs) and the Compressed Sparse Row (CSR) matrix format.

Index Terms—Sparse Matrix, SpMV, CUDA, Parallelization, Storage Format

I. INTRODUCTION

Sparse Matrix-Vector Multiplication involves multiplying a sparse matrix in which most elements are zero, by a dense vector. Due to its low computational intensity and irregular memory access patterns, SpMV is often **memory-bound**, making its efficient execution crucial for the overall performance of many high-performance computing applications.

Parallelizing SpMV poses significant challenges, mainly for the structure of the sparse matrix itself, which leads to unpredictable memory access and poor cache utilization. Unlike dense matrix operations, the sparsity pattern is not uniform and can vary widely across applications, making load balancing across threads or processing units difficult.

Moreover, implementing SpMV into modern platforms requires consideration of **hardware-specific features** like warp-based execution, memory coalescing, or limited on-chip memory. As a result, multiple storage formats and algorithmic variants have been developed to adapt SpMV to specific hardware characteristics.

Despite its conceptual simplicity, SpMV remains a challenging and active area of research in various research fields.

II. PROBLEM STATEMENT

Sparse Matrix-Vector multiplication is a well-known problem whose purpose is to calculate the product between a sparse matrix, that is, whose values are mostly zeros and therefore do not contribute to the final result, and a dense vector. The product between a matrix A of size $n \cdot m$ and a vector x (which **must be** of size m) gives as a result a vector y of size n that can be calculated given the following formula:

$$y_i = \sum_{j=0}^m A_{ij}x_j$$

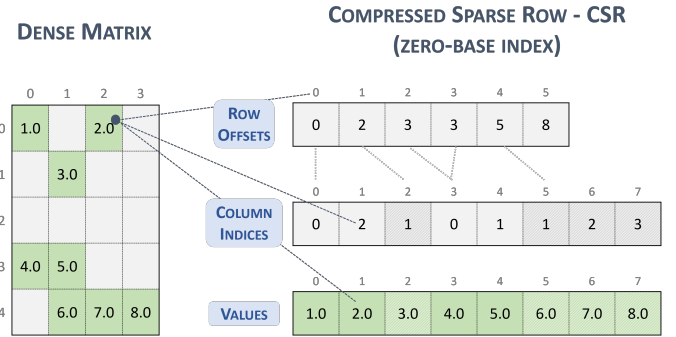
Where y_i is the i -th element of the resulting vector.

III. STORAGE FORMAT

Storing a large sparse matrix as a dense one is not memory efficient at all since most of the values are zeros and therefore do not contribute to the final result.

To store sparse matrices efficiently many storage formats were, and still are, being developed such as the **Compressed Sparse Row** matrix format, CSR in short.

Fig. 1. CSR matrix format



This format consists of three arrays used to store the rows, columns and non-zero values of the matrix, the main advantage is that the array containing the row indices is constructed by sorting and creating a prefix sum of the original row indices reducing the length of the row array from the number of non-zero to the number of rows plus one.

Using the CSR format, access to the i -th row can be obtained just by taking the i -th value of the row array as shown in figure 1.

IV. PARALLELIZATION

Starting from a full CPU implementation as baseline, four different CUDA implementation were developed, each one being a variation of the previous one taking into account a specific group of operations.

A. One Thread Per Row

For the first CUDA implementation, since each value of the output array can be calculated independently from the others, the naive approach is to create a single thread for each row which multiplies all its non-zero values with the corresponding input vector elements and sums them together sequentially.

B. Additional Sorting Kernel

Since the CSR format requires sorting, for large matrices a lot of time is spent for the construction of the row array of the matrix so as a solution a **Bitonic Parallel Sort** algorithm was implemented to achieve faster execution time.

C. One Thread Per Non-Zero

For matrices having long rows the naive implementation can be slow so as an alternative approach the matrix can be subdivided into chunks having one thread for each non-zero value.

The main difficulties of this approach are the subdivision of the matrix into groups to balance the workload and the additional **concurrency** problem caused by access to the same memory region between multiple threads.

As a naive solution for the concurrency problem, mainly caused by the sum of the products, each product is stored inside the matrix itself and the sum is executed sequentially on the CPU.

D. Warp Reduction

A better solution to the previous problem is to exploit modern warp **reduction primitives** to parallelize the sum operations hence increasing performance of the algorithm.

V. STATE OF THE ART

From the original Compressed Sparse Row (CSR) [1] implementation many more research has been conducted to improve its use in algorithms such as Sparse Matrix-Vector Multiplication.

One of the main problems of the format was its performance which is **highly dependent** on the used hardware and the actual structure of the matrix. To achieve this, several optimized variants of CSR have been developed such as CSR5 and its extension CSR5BC which improve parallelism and memory access on modern GPUs using segmented sum techniques and index compression [2].

The CSR&RV format reduces memory usage by storing redundant values only once, significantly lowering storage costs and boosting throughput [3].

On heterogeneous systems, formats like CSR-k restructure data hierarchically for better adaptability across CPUs and GPUs [4]. Auto-tuning frameworks such as AlphaSparse generate customized kernels tailored to both sparsity patterns and hardware, delivering further gains [5].

Additionally, emerging architectures like Processing-In-Memory (PIM) are prompting new SpMV implementations, such as SparseP, that harness in-memory computing advantages [6]. Together, these innovations demonstrate that while CSR remains foundational, ongoing adaptations are essential for maximizing SpMV efficiency on evolving platforms.

VI. METHODOLOGY AND CONTRIBUTIONS

For profiling on the CPU the `gettimeofday` function was used to measure the time taken by different operations down

to the microseconds, such as memory allocation, matrix file parsing, input vector value generation, etc....

For the GPU **CUDA events** were used instead to achieve **better accuracy** for the elapsed time including the internal synchronization of the threads.

The naive GPU kernel implementation is expected to be **slower** for long row matrices since the SpMV becomes almost sequential, which can be seen in the following pseudo-code:

Algorithm 1 Naive GPU SpMV kernel

Require: The input matrix A of size $n \cdot m$ and the vector x of size m .

```

1: procedure SPMV( $A, x, n, m$ )
2:    $y \leftarrow \emptyset$ 
3:   for  $r$  in  $\{1 \dots n\}$  parallel do
4:     for  $c$  in  $\{1 \dots m\}$  do
5:        $y[r] + = A[r \cdot m + c] + b[c]$ 
6:     end for
7:   end for
8:   return  $y$  ▷ the result vector
9: end procedure

```

The other approach using one thread for each non-zero element is expected to be more **consistent** across different matrices and slightly faster than its naive counterpart since it is also able to parallelize the sums even if the implemented algorithm can still be improved.

Algorithm 2 Warp reduction SpMV kernel

Require: The input matrix A of size $n \cdot m$ and the vector x of size m .

```

1: procedure SPMV( $A, x, n, m$ )
2:    $y \leftarrow \emptyset$ 
3:   for  $i$  in  $\{1 \dots n \cdot m\}$  parallel do ▷ product only
4:      $c \leftarrow cols[i]$ 
5:      $A[i] * = x[c]$ 
6:   end for
7:   for  $warp$  in  $\{1 \dots n \cdot m / 32\}$  parallel do ▷ warp reduction
8:      $val \leftarrow A[warp \cdot 32]$ 
9:     for  $off$  in  $\{16, 8, 4, 2, 1\}$  do
10:      for  $j$  in  $\{1 \dots 32\}$  parallel do
11:         $val + = A[warp \cdot 32 + j]$ 
12:      end for
13:    end for
14:     $A[warp \cdot 32] = val$ 
15:  end for
16:  for  $r$  in  $\{1 \dots n\}$  do ▷ sum of the reductions
17:    for  $c$  in  $\{1 \dots m / 32\}$  do
18:       $y[r] + = A[r \cdot m + c \cdot 32]$ 
19:    end for
20:  end for
21:  return  $y$  ▷ the result vector
22: end procedure

```

VII. SYSTEM DESCRIPTION AND EXPERIMENTAL SET-UP

A. System Description

For the development of the project the following requirements were used inside the university cluster:

- SLURM 32.11.3
- CUDA 12.1
- GCC 12.3.0
- GNU Make 4.3
- Eigen 12.3.0 (for the utilities)
- Python 3.9.18 (for the utilities)

The project can be compiled using the provided Makefiles with GCC and CUDA compilers and can be run into the cluster via the SLURM workload manager utilities.

TABLE I
SYSTEM DETAILS

System	Processor	Cores per Socket	RAM	Accelerator
baldo	Intel(R) Xeon(R) Silver 4214 CPU	12 at 2.2 GHz	405 GB	
edu-short	AMD EPYC 9334 32-Core Processor	32 at 3.9 GHz	810 GB	NVIDIA L40S

B. Dataset description

Eight different matrices were selected as dataset, each with specific characteristics to evaluate the performance of the various implementations.

The first five matrices are mostly symmetric and differ primarily in the number of non-zero elements, which increases progressively from one matrix to the next.

The `heart1` matrix stands out for its **higher density**, with non-zero values accounting to approximately 10% of the total entries.

The final two matrices are asymmetric and are designed to stress test different dimensions: one has a large number of rows, while the other has a large number of columns.

TABLE II
MATRIX DETAILS

Name	Rows	Columns	Non-Zeros	Symmetric
1138_bus	1,138	1,138	4,054	Yes
bcstk15	3,948	3,948	117,816	Yes
bcstk31	35,588	35,588	1,181,416	Yes
pwtk	217,918	217,918	11,524,431	Yes
cage15	5,154,859	5,154,859	99,199,551	No
heart1	3,557	3,557	1,385,317	No
relat8	345,688	12,347	1,334,038	No
connectus	512	394,792	1,127,525	No

C. Experimental Set-up

For the experimental setup, each main operation is profiled by measuring its **execution time**.

The profiled operations include:

- 1) Memory allocation and copy
- 2) File Input/Output
- 3) Sorting for CSR matrix construction
- 4) Matrix-Vector multiplication

The CSR matrix is constructed only once at the start of the program. To ensure accurate timing of the matrix-vector multiplication, four initial **warm-up cycles** are executed and excluded from profiling. At the end the average value between 10 different runs is evaluated.

VIII. EXPERIMENTAL RESULTS


The results for the different implementations across all matrices are shown in the following graphs.

A few important points should be noted:

- Results for the naive CPU and GPU implementations on larger matrices are missing. This is because their total execution time exceeded the 5-minute time limit enforced by the cluster;
- The Y-axis of the graph uses a **logarithmic scale** rather than a linear one.

The logarithmic scale is necessary because the first five matrices increase in size by roughly an order of magnitude each, leading to similarly scaled increases in execution times.

Fig. 2. Full Time Benchmark



From the results, we can observe that the CPU implementation performs better on smaller matrices, primarily due to the **lower overhead** compared to the GPU (e.g. memory allocation and kernel dispatch). However, its performance degrades rapidly as the number of non-zero elements increases.


Since sorting the matrix rows for CSR construction takes a significant portion of the total time, using a Bitonic Parallel Sort considerably reduces overall execution time.

Other **performance-impacting factors** include memory allocation and data transfer, especially between host and device, as well as file I/O operations. These are areas that could benefit from further optimization.

As expected, using one thread per row yields better performance on matrices with a large number of rows (e.g. `relat8`). However, it performs poorly on matrices with fewer but longer rows (e.g. `connectus`).

On the other hand, using one thread per non-zero element delivers more consistent performance across different matrix shapes. Nevertheless, since this approach requires the summation of partial products to be performed on the CPU (to avoid race conditions due to concurrent memory access), it is often slower than the one-thread-per-row implementation.

Fig. 3. SpMV Time Benchmark



The warp reduction strategy, which partially parallelizes the summation process, shows **limited performance improvement**. This is expected, as only one full reduction is performed on the device for each warp, while the remaining summation still occurs on the host. This bottleneck is especially evident in the case of the `relat8` matrix, where the high number of rows leads to more CPU-side operations.

IX. CONCLUSION

To summarize, parallelizing both the sorting phase and the matrix-vector product on the GPU significantly improves performance for large sparse matrix-vector multiplications (SpMV).

While the one-thread-per-row approach generally delivers better performance, the one-thread-per-non-zero implementation has potential for textbfurther optimization and could be improved to achieve lower execution times.

Looking ahead, future work could explore **adaptive** or hybrid algorithms that dynamically select the most suitable strategy based on the specific characteristics of each matrix, as well as different approaches to balance the parallel operations across all the elements of the matrix.

REFERENCES

- [1] S. C. Eisenstat, M. C. Gursky, M. H. Schultz, and A. H. Sherman, *Yale Sparse Matrix Package. II. The Nonsymmetric Codes*, Aug. 1977. [Online]. Available: <https://doi.org/10.21236/ada047725>
- [2] W. Zhang, Y. Chen, B. Zhang, X. Shen, J. Xue, and C. J. Xue, "Csr5: An efficient storage format for cross-platform sparse matrix-vector multiplication," in *Proceedings of the 29th International Workshop on Languages and Compilers for Parallel Computing (LCPC)*. Springer, 2016, pp. 77–93.
- [3] Z. Tang, X. Zhu, L. Wang, W. Song, D. Cheng, and Q. Wang, "Csr&rv: A compressed sparse row format with reduced value redundancy," in *Proceedings of the 19th International Conference on Algorithms and Architectures for Parallel Processing (ICA3PP)*. Springer, 2023.
- [4] W. Liu, J. Chen, G. Lin, X. Wang, and M. Liu, "Csr-k: A fast and portable spmv kernel with hierarchical row blocking," *arXiv preprint arXiv:2203.05096*, 2022. [Online]. Available: <https://arxiv.org/abs/2203.05096>
- [5] J. Yin, X. Li, Y. Li, X. Wang, and J. Chen, "Alphasparse: Auto-tuning and accelerating spmv with machine-designed formats," *arXiv preprint arXiv:2212.10432*, 2022. [Online]. Available: <https://arxiv.org/abs/2212.10432>

- [6] S. Lee, M. Kim, M. Joung, J. Park, and J. W. Lee, "Sparsep: Toward realistic performance modeling of sparse matrix-vector multiplication on processing-in-memory systems," *Proceedings of the ACM on Measurement and Analysis of Computing Systems (POMACS)*, vol. 6, no. 1, pp. 1–35, 2022.